

PEXpo-Bench: validate the execution environment before trusting LLM architecture effects


1,027 curated personal-exposure items

5 sub-domains · 3 question types
gold answers independently re-verified



5 complete configurations × 4 LLMs

Weakest model (GPT-5.4-nano)



Corrected rerun, all four models

